

Course: Artificial Neural Networks and Deep Learning

Exercises Unit 2: Deep Learning Methods

This set of exercises has been designed to experiment with different hyperparameter values in order to analyze how they affect the performance of a neural network. The following files are used:

- Median_house_value.ipynb
- MedianHouseValuePreparedCleanAttributes.csv
- MedianHouseValueOneHotEncodedClasses.csv

Exercise 1. Performance evaluation

Determine the following values: train error, test error (using dev set), bias, variance and training time. Repeat the execution at least three times. Consider 5% as Bayesian error (human error).

Exercise 2: Changing basic hyperparameters

Change hyperparameters related to: batch size, number of layers, and number of neurons. Estimate train error, test error, bias, variance and training time. Consider 5% as Bayesian error (human error).

Exercise 3: Changing activation functions

Change the activation function in the following part of the code:

```
for neurons in n_neurons_per_hlayer:
    model.add(keras.layers.Dense(neurons, activation="relu"))
```

Possible activation functions: elu, relu, tanh, linear

(Other functions: sigmoid, softsign, softplus, hard_sigmoid, exponential, softmax, selu)

More information about activation functions: <https://conx.readthedocs.io/en/latest/ActivationFunctions.html>

Exercise 4: Changing initializers

Change the lines:

```
my_initializer= keras.initializers.RandomUniform(minval=-0.05, maxval=0.05, seed=None)
for neurons in n_neurons_per_hidden_layer:
    model.add(keras.layers.Dense(neurons, activation="relu", kernel_initializer=my_initializer))
```

Possible initializers:

```
keras.initializers.Zeros()
keras.initializers.RandomUniform(minval=-0.05, maxval=0.05, seed=None)
keras.initializers.RandomNormal(mean=0.0, stddev=0.05, seed=None)
keras.initializers.he_uniform(seed=None)
keras.initializers.he_normal(seed=None)
```

See this link for more information about initializers: <https://keras.io/initializers/>

Exercise 5: Using batch normalization

Change the lines where the model is created.

Normalize after activation function:

```
model = keras.models.Sequential([
    keras.layers.InputLayer(input_shape=(INPUTS,)), batch_size=None),
    keras.layers.Dense(500, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(250, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(75, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(25, activation="elu", kernel_initializer="he_normal"),
    keras.layers.BatchNormalization(),
    keras.layers.Dense(3, activation="softmax")])
```

Normalize before activation function:

```
model = keras.models.Sequential([
    keras.layers.InputLayer(input_shape=(INPUTS,)), batch_size=None),
    keras.layers.Dense(500, kernel_initializer="he_normal", use_bias=False),
    keras.layers.BatchNormalization(),
    keras.layers.Activation("elu"),
    keras.layers.Dense(250, kernel_initializer="he_normal", use_bias=False),
    keras.layers.BatchNormalization(),
    keras.layers.Activation("elu"),
    keras.layers.Dense(75, kernel_initializer="he_normal", use_bias=False),
    keras.layers.BatchNormalization(),
    keras.layers.Activation("elu"),
    keras.layers.Dense(25, kernel_initializer="he_normal", use_bias=False),
    keras.layers.BatchNormalization(),
    keras.layers.Activation("elu"),
    keras.layers.Dense(3, activation="softmax")])
```

Exercise 6: Using regularization

In this exercise, analyze two different cases: (a) batch size=512 and (b) batch size=256.

PART (1): Use regularizer L2 as follows:

```
model = keras.Sequential(name="my_model")
model.add(keras.layers.InputLayer(input_shape=(INPUTS,)), batch_size=None))
for neurons in n_neurons_per_hidden_layer:
    model.add(keras.layers.Dense(neurons, activation="relu",
                                  kernel_regularizer=keras.regularizers.l2(0.001)))
model.add(keras.layers.Dense(OUTPUTS, activation="softmax",
                              kernel_regularizer=keras.regularizers.l2(0.001)))
```

Possible regularizers:

```
keras.regularizers.l1(0.001)
keras.regularizers.l2(0.001)
```

Analyze regularizers with different values for lambda (e.g., 0.001, 0.0001, ...)

PART (2): Use dropout as follows:

```
model = keras.models.Sequential([
    keras.layers.InputLayer(input_shape=(INPUTS,)), batch_size=None),
    keras.layers.Dense(500, activation="elu", kernel_initializer="he_normal"),
    keras.layers.Dropout(rate=0.01),
    keras.layers.Dense(250, activation="elu", kernel_initializer="he_normal"),
    keras.layers.Dropout(rate=0.01),
    keras.layers.Dense(75, activation="elu", kernel_initializer="he_normal"),
    keras.layers.Dropout(rate=0.01),
    keras.layers.Dense(25, activation="elu", kernel_initializer="he_normal"),
    keras.layers.Dropout(rate=0.01),
    keras.layers.Dense(3, activation="softmax")])
```

Analyze different values of dropout rate (e.g., rate = 0.1, 0.01, ...)

Exercise 7: Changing learning rate and epochs

Change learning rate and number of epochs. Their values are written in the following lines of code:

```
n_epochs = 1000
```

```
learning_rate = 0.1
```

Exercise 8: Changing optimizers

Change optimizer:

```
model.compile(loss=tf.keras.losses.categorical_crossentropy,
               optimizer=tf.keras.optimizers.SGD(learning_rate=learning_rate),
               metrics=["categorical_accuracy"])
```

Optimizers:

```
tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.9)
tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.9, nesterov=True)
tf.keras.optimizers.RMSprop(learning_rate=learning_rate, rho=0.9)
tf.keras.optimizers.Adam(learning_rate=learning_rate, beta_1=0.9, beta_2=0.999)
```

NOTE: Try also smaller learning rates for RMSprop and Adam (0.01, 0.001, etc.)

Other optimizers: https://www.tensorflow.org/api_docs/python/tf/keras/optimizers

Exercise 9: Multiple changes

Combine multiple changes of hyperparameters to obtain the best performance.